

Enriching Exploratory Spatial Data Analysis with modern computer tools

Robin CURA, PhD Student in Geography

Univ. Paris 1 Panthéon-Sorbonne & Géographie-cités laboratory

 **ECTQG 2019**

Handling massive data sources with classical tools

An intermediate kind of data

Handling the recent massive data sources

~~"Big Data"~~

- Data that can't fit on a personal computer
- Few evolutions in the last years
- Still the usual distributed computing solutions : Hadoop, Apache Spark
- Still hard to grasp for non computer scientists

Traditionnal ~~large data~~

e.g. census data, territorial mesh data etc.

- Can be managed in most analysis softwares
- Often adapted for spreadsheets, GIS, GUI softwares (GeoDA...) etc.
- Can also be analyzed with CLI analysis tools like R or Python (cf. yesterday's workshop)

What lies in-between these dataset types ?

An intermediate kind of data

Data			Storage and analysis	
Quantity	Memory size	Examples	Storage infrastructure	Analysis and visualization tools
Up to 1,000 rows	~1 MB	Aggregated census data	Text file	Spreadsheets, GIS
1,000-100,000 rows	~1-50 MB	Detailed Census data	Text files	GIS, GUI software (GeoDA, Tableau...) etc.
100,000 – a few million rows	~50 MB – 1 GB	Individual data, time series of multiple sensors	Spreadsheet or binary files (SHP, geopackage...)	Interactive (GUI) statistical tool or Command-Line Interface (CLI) tools
10 to 100 million rows	~1 - 10 GB	Many new open datasets : equipments, user-generated content, VGI etc.	???	???
> 100 million rows	> 10s of GB	Spatio-temporal data, automated reporting, big companies datasets ...	Distributed Databases	High-Performance Computing

Handling massive data

It's not really a new problem


- Traditionnal handling : relational DMBS, e.g. MySQL, PostgreSQL, SQLite
 - Great for :
 - archiving large data
 - multiple users and concurrent queries
 - updatability
 - diversity of types and queries
 - customisation
 - universality through SQL
 - Issues with : **speed**
 - install/setup : can't be setup in a few minutes
 - data import : made for updating, slow for massive imports
 - queries : row-based DMBS : slow for global data aggregates/joins

Traditionnal relationnal DBMS are not that good a fit for Exploratory Data Analysis (EDA).

- Mostly a single-user context, requiring many back and forth between global structure and peculiarities

A few benchmarks

Performance of various DataBase Management Systems



R. Cura (2018), updated from S. Pafka (2017)
* Omnisci (CPU) and SQLite : Benchmark executed on a less powerfull computer

Using the new generation of colum-based DBMS

Advantages :

- Easy to setup
- Fast for data insert
- Fast (extremely) for column-based operations : agregation, joins
- Can be queried through SQL (at least some of them)

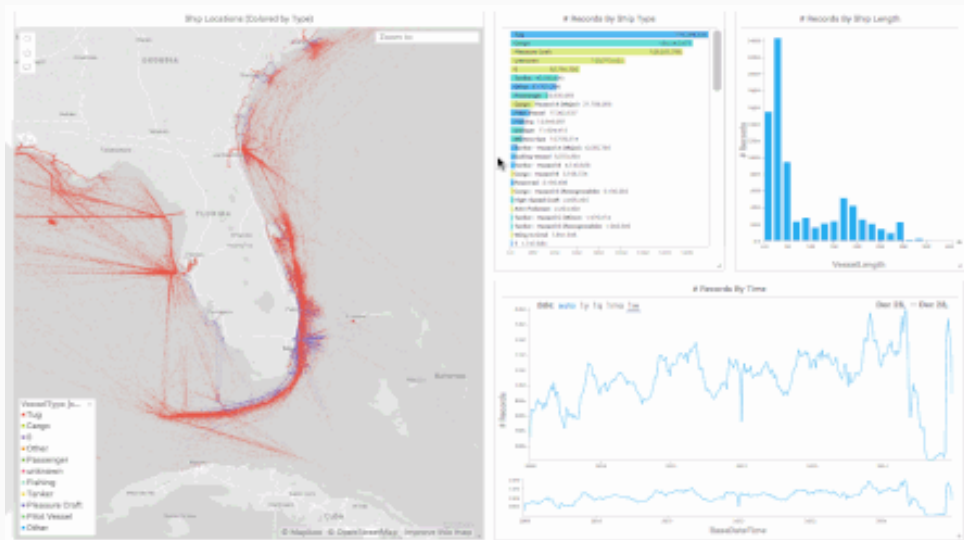
Weaknesses :

- Not performant for upserts and data updates
- Not as mature as relational DBMS : still the state of the art, although used a lot by big tech companies
- Less operators/flexibility on queries (especially spatial queries : nothing can reach PostGIS spatial queries)

An example of a columnar DBMS

Presenting the example of Omnisci (former MapD), an open-source* columnar DBMS

- A DBMS made for GPU processing :
 - Can query hundreds of billions of rows in a blink
 - When it is run on an adapted hardware configurations : high-end GPUs etc.
 - Check www.omnisci.com/demos :



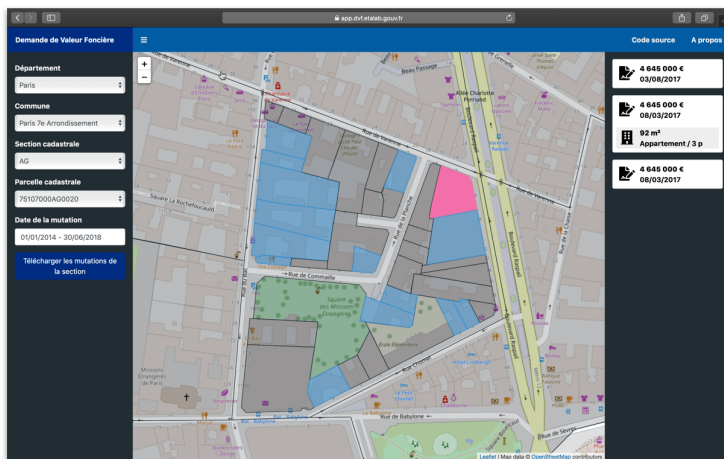
- Still works very well for 10^6 - 10^8 rows-data on an aging personal computer without GPU

Illustrating an EDA approach using Omnisci

Example dataset :

The recent government opendataset on real estate transactions : the DVF ("Demande de Valeurs Foncières")

- Logs all (public and private) real estate transactions on *almost* all of France
- Detailed geolocation : cadastral parcel (e.g. big as a building)
- Detailed price for each transaction, with the corresponding informations on the real estate types, area and composition
- Yearly, since 2014



- 1 year : ~ 3M rows : easy to manage through CLI
- 5 years (2014-2018) : ~15M rows, doesn't fit in memory
- A good candidate for a quick test/demo

Illustrating an EDA approach using Omnisci

Preparing the DBMS :

- Run through Docker container : `docker run -name testdb -d -v $HOME/omnisci-docker-storage:/omnisci-storage -p 6273-6280:6273-6280 omnisci/core-os-cpu`
- Prepare the database :
 - Log to Omnisci container : `docker container exec -ti testdb /bin/bash`
 - Run the sql shell and create the table :

```
CREATE TABLE dvf ( id_mutation TEXT, annee SMALLINT, [ ... ] );
```

- Populate the table : `COPY dvf FROM '/data/dvf_2014-2018.csv.gz';` (~couple minutes)
- Verify that the data looks correct (sub-seconds query)

```
omnisql> SELECT COUNT(*) AS n_rows FROM dvf;  
> n_rows  
> 13 255 975
```

The full DB infrastructure can be setup in a few lines and ready in a few minutes

Exploring the database with R

Omnisci has a SQL interface, through ODBC/JDBC

-> It can be interactively queried from both R (using `RJDBC`) and from Python (`pymapd`)

-> Here, examples using R and tidyverse-style piped queries through the common database connection `DBI`

```
# Connecting to the DB
```

```
db_connection <- DBI::dbConnect(drv = db_driver, [credentials],  
                               url = "jdbc:omnisci:localhost:6274:omnisci")
```

```
# Loading the table
```

```
dvf_data <- tbl(src = db_connection, "dvf")
```

```
dvf_data
```

```
## # Source:   table<dvf> [?? x 22]  
## # Database: JDBCConnection  
##   id_mutation annee date_mutation dept  code_commune section id_parcelle nature_mutation valeur_fonciere nombre_lots type_local surface_reelle_  
##   <chr>       <dbl> <chr>      <chr> <chr>      <chr>      <chr>      <chr>      <chr>      <dbl>      <dbl> <chr>      <dbl>  
## 1 2018-1      2018 2018-01-03 01    01053      010530... 01053000AN... Vente      109000      1 Dépendance      NA  
## 2 2018-1      2018 2018-01-03 01    01053      010530... 01053000AN... Vente      109000      2 Apparteme...    73  
## 3 2018-2      2018 2018-01-04 01    01095      010950... 01095000AH... Vente      239300      0 Maison          163  
## 4 2018-2      2018 2018-01-04 01    01095      010950... 01095000AH... Vente      239300      0 Maison          51  
## 5 2018-2      2018 2018-01-04 01    01095      010950... 01095000AH... Vente      239300      0 Maison          51  
## 6 2018-2      2018 2018-01-04 01    01095      010950... 01095000AH... Vente      239300      0 Maison          163  
## 7 2018-3      2018 2018-01-04 01    01343      013430... 01343000ZR... Vente      90000      0 NA              NA  
## 8 2018-3      2018 2018-01-04 01    01343      013430... 01343000ZR... Vente      90000      0 Maison          150  
## 9 2018-3      2018 2018-01-04 01    01343      013430... 01343000ZR... Vente      90000      0 NA              NA  
## 10 2018-6     2018 2018-01-04 01    01053      010530... 01053000BD... Vente      67000      1 Apparteme...    45  
## # ... with more rows, and 10 more variables: nombre_pieces_principales <dbl>, nature_culture <chr>, nature_culture_speciale <chr>,  
## #   surface_terrain <dbl>, latitude <dbl>, longitude <dbl>, surface2 <dbl>, type_surface <chr>, prix_surface <dbl>, mois <dbl>
```

Exploring the database with R

Basic EDA

```
dvf_data %>% group_by(annee, mois) %>%  
  summarise(nb_transactions = n(), meanPrice = mean(valeur_fonciere, na.rm = TRUE)) %>%  
  ungroup() %>% arrange(annee, mois) %>%  
  collect() %>% # Run computations inside the DB , retrieve the results locally  
  gather(Var, Value, -annee, -mois) %>%  
  ggplot() + aes(mois, Value, fill = mois) + geom_col() +  
  facet_grid(Var~annee, scales = "free_y")
```

Mean transaction price and number of transaction through months and years



Handling massive data sources with classical tools

- Many new-generations DBMS (NoSQL, graph DB, collection/document DB etc.)
- Among these, the relational column-based DBMS can offer a known and almost-universal interface (SQL) and integrate very easily our already existing workflows for EDA
- This allows to scale up, for a few order of magnitudes, the amount of data that *any* quantitativist geographer can now analyse
- Omnisci (but likewise Yandex's ClickHouse, MonetDB, Amazon Redshift, DuckDB, Uber's AresDB (soon) etc.) offers a quick-to-setup interface to managing such data

What about spatial data visualisation ?

What about spatial data visualisation ?

Often requires visualization of the spatial data

- Huge historical strength of GIS
 - Very good interaction with DBMS (e.g. QGIS was conceived as a PostGIS viewer)
 - Still disruptive inside an EDA CLI-based workflow
- CLI (Python/R) :
 - Many static visualisations : plot(SPDF), matplotlib, geoPandas.plot, ggplot2(sf) etc.
 - Yet, Shneiderman's mantra : *"Overview First, Zoom and Filter, Then Details-on-Demand"*
- For a few years, domination of Leaflet :
 - Vectorial data allows powerfull interactions
 - Limited in data quantities
 - Use of raster- (or vector-) tiles to handle larger data
 - Lacks of interactivity

What about spatial data visualisation ?

Main problem : Leaflet can't support CLI-size data :


- Hard to display/interact with more than a few hundreds polygons (~ a few thousands points)

A very simple example of the limits :

- The DVF dataset cannot be represented at France's scale, only at the "département" level
- To add simple spatial visualisations of the DB-stored DVF dataset, we would need to use a GIS...

At the same time

- Many new technologies allow for webGUI-based analysis : MapboxGL, DeckGL, KeplerGL etc...
- As much / more performant than GIS or any local software



What about spatial data visualisation ?

Both R and Python can use performant visualisation solutions :

- PyOpenGL
- PyViz stack (Bokeh, Holoviews, hvPlot...)
- Plotly
- mapboxgl-jupyter
- RGL
- Plotly
- mapdeck
- leafgl

Focus on leafgl (R)

- The versatility of leaflet (a wrapper around the Leaflet.glify JS library)
- Ability to interact with thousands of spatial entities
- Without leaving the CLI environment



What about spatial data visualisation ?

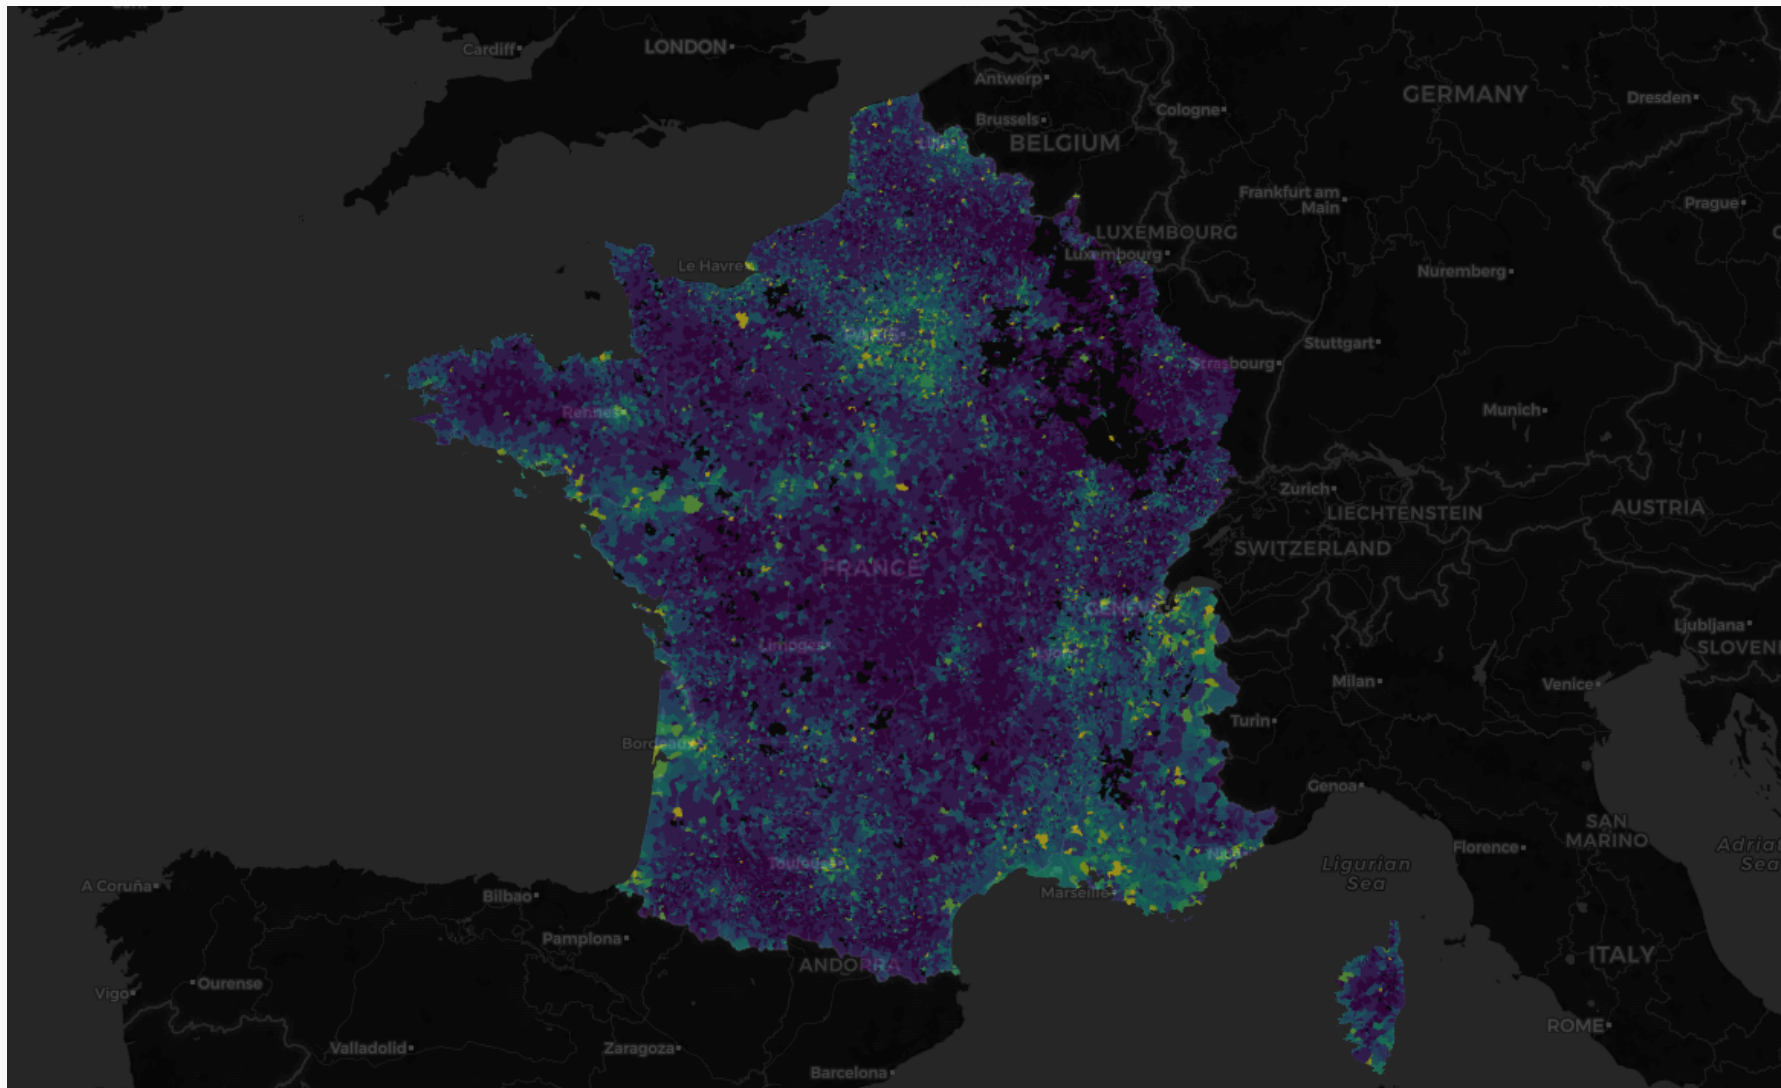
An example with our DVF dataset :

- Aggregation on the ~36000 "communes"-scale
- Computing the mean square-meter price of the transaction for all years

N.B. : it can also be integrated inside more complete and ad-hoc applications (Shiny, Dash, Bokeh...)

```
# [ ... ] Preparing the data and loading the usual spatial packages
# [ ... ] Discretizing the values
# Making the map
leaflet() %>%
  addProviderTiles(provider = providers$CartoDB.DarkMatter) %>%
  addGlPolygons(data = map_data, color = communes_colors) # replaces leaflet::addPolygons
```

Example



Conclusion :

- Intermediate solutions for intermediate data quantities exist
- Easy to setup/understand, doesn't require much computer science skills or computing power
- Just one small step ahead what you saw at yesterday's workshop
- Don't get blocked by data that are "a bit too big"